

Debanjan Mazumdar

AI Product Manager | GenAI • Multi-Agent Systems • LLM • RAG • Data Products

debanjan186@gmail.com • 0755807658 • Paris, France • [linkedin.com/in/dmazumdar](https://www.linkedin.com/in/dmazumdar) • github.com/dmazumdar186

AI Product Manager with 15 years of experience in data-intensive environments, specialized in the design and deployment of production-grade Generative AI capabilities (LLM, RAG, Multi-Agent Systems). Expert at scoping complex AI features — from discovery to production — with a strong grasp of ethical, regulatory (GDPR) and governance constraints, paired with a hands-on builder practice (Python, Cloudflare Workers, Modal cron, self-iterating audit loops).

Results: +30 % adoption • +20 % CSAT • -40 % latency • -35 % iteration cycles • +40 % BU adoption

PROFESSIONAL EXPERIENCE

AI Product Manager

Wiser Solutions, Paris | Nov. 2022 – Present

- Designed and shipped production-grade GenAI capabilities (triage, recommendation, RAG customer support, smart alerts) using OpenAI Assistants, Claude and multi-agent systems (MCP, A2A) → **-40 % latency, +25 % precision, +25 % adoption**
- Defined the AI vision and capability roadmap with value hypotheses, go/no-go criteria and rollback plans; aligned **5 cross-BU teams** (Data/AI, MLOps, architecture, security, business) on a shared vision
- Authored AI-oriented PRDs (API/data contracts, evaluation thresholds, non-determinism handling, GDPR/privacy-by-design) and led technical feasibility reviews with Data/AI teams
- Orchestrated the global GTM rollout with enablement plans and field playbooks → **+40 % BU adoption, +30 % user adoption, +20 % CSAT**
- Built analytics dashboards (usage / coverage / quality) and operational signals (drift, incidents) → **+25–30 % delivery precision, -25 % sprint ambiguity** via DoR/DoD and pre-prod QA

AI Product Engineer — Freelance Engagement

Accessory Masters (Elite Broker Group), remote | Dec. 2025 – Mar. 2026

- Designed and shipped an autonomous cold-email pipeline at scale (**24,000 emails/month, ~800/day across 32 warmed inboxes**) with Claude-powered personalization, reply classification and contextual auto-reply → **4 % reply rate**, 3-hour hot-lead SLA via Telegram
- Architected a serverless event-driven stack (Cloudflare Workers + KV cron, Instantly.ai webhooks, GoHighLevel CRM V2 API, scheduled Modal jobs) with idempotency keys, 60-day KV dedup TTL, and LLM guardrails (limits, voice, fallbacks)
- Productized monitoring & governance: operator dashboard (Vercel), **/api/health** endpoint, **--dry-run** modes on every paid path, synthetic canary → silent drift caught before any API-credit impact

Data Product Manager

InfoTnT, Paris | Jun. 2021 – Nov. 2022

- Led product discovery for a data-driven recommendation engine (quant/qual, JTBD) → **-35 % iteration cycles, +25 % post-launch fit**
- Translated business needs into backend capabilities, epics and data contracts with success metrics and integration constraints
- Standardized prompts, templates and versioning procedures for AI components → predictable behaviors and reliable releases; coordinated releases via PRDs/RFCs and functional QA

Senior Data Product Owner

Pitney Bowes Inc, Pune | Apr. – Sep. 2019

- Reduced time-to-market by **-20 %** via cross-squad dependency mapping, tightened release cadence, pre-prod checklists and structured RAID reviews

Senior Data Product Owner

Evolent International, Pune | Jun. 2018 – Feb. 2019

- Improved platform scalability and performance (**+30 %**) by introducing SLA/SLO to arbitrate run/change priorities; embedded governance and QA checkpoints into the delivery flow

Senior Product Owner

Avaya India Pvt Ltd, Pune | Jul. 2015 – Mar. 2018

- Accelerated delivery velocity (+30 %) through stronger Scrum discipline and OKR ↔ roadmap alignment; reduced requirement instability (–25 %) to improve release readiness

QA Engineer / Release Coordinator — IDrive India Pvt Ltd, Bengaluru | Apr. 2013 – Jul. 2015

Software Engineer — Tata Consultancy Services, Bengaluru | Nov. 2010 – Mar. 2013

SKILLS

AI & GenAI :	LLM, RAG, Agentic AI, Multi-Agent Systems, OpenAI Assistants, Claude, MCP, A2A, supervised / unsupervised ML, AI evaluation frameworks (precision, coherence, robustness), non-determinism handling, prompt caching, LLM guardrails, self-iterating audit loops (anneal), model routing (OpenRouter), Cloudflare Workers, Modal cron, Firecrawl
Product Mgmt :	Vision & Roadmap, Discovery (quant/qual, JTBD), PRD, API/data contracts, Backlog, KPI/OKR, A/B testing, DoR/DoD, Enablement & GTM
Governance :	GDPR / privacy-by-design, access control, audit trails, tech sovereignty, AI ethics
Collaboration :	Cross-functional (Data/AI, Dev, CSM, Sales, MLOps, Security), workshop facilitation, stakeholder management
Tools :	Jira, Confluence, Google AntiGravity, N8N, Make, Figma, Miro, Mixpanel, SQL, GenAI APIs (OpenAI, Google, Anthropic)

EDUCATION

MSc International Strategic Business

Toulouse Business School, Paris | 2019–2021

BE Information Technology Engineering

CMR Institute of Technology, Bengaluru | 2006–2010

LANGUAGES

English : Bilingual • French : Bilingual • Hindi : Native • Bengali : Native

PERSONAL PROJECTS

- **Anneal** (Mar. 2026) — Python CLI for self-iterating audit loops over git diff (classic + adversarial Red-vs-Blue); dual-adapter LLM (Anthropic + OpenRouter), cheap / balanced / premium tiers; terminates on 2 consecutive zero-bug rounds
- **YouTube Video Analyzer** (May 2026) — frame-by-frame breakdown of YouTube videos: PySceneDetect + perceptual dedup + 3×3 tiling (–85 % vision tokens) + multi-model routing (Claude / Gemini free / OpenRouter); 73 tests, 8 clean audit rounds
- **Job Tracker PM France** (May 2026) — daily ETL pipeline: 5 job boards scraped (WTTJ, Indeed, APEC, France Travail, Google Jobs), ICP filtering, contact enrichment via INSEE SIRENE + Firecrawl, HTML digest by email; stack Firecrawl, Serper, Modal cron, SQLite
- **Self-Outbound Engine** (Jan. 2026) — autonomous prospecting engine: 180–270 emails/day across 6–9 warmed inboxes, reply classification, Cal.com auto-reply with human-randomized delay, Telegram alerts; stack Cloudflare Workers, Apollo, Gemini, Claude Haiku
- **CV Optimizer Agent** (Apr. 2026) — AI agent (Streamlit + Gemini): ATS scoring of CV vs JD, generation of optimized CV & cover letter PDFs, multilingual
- **ProdCraft** (YouTube, Sep. 2025 – Present) — ed-tech channel for aspiring Product Managers: fundamentals & best practices