

Debanjan Mazumdar

AI Product Manager & Builder | GenAI • RAG • Agentic AI • LLM • Automation — shipped to production

debanjan186@gmail.com • 0755807658 • Paris, France • [linkedin.com/in/dmazumdar](https://www.linkedin.com/in/dmazumdar) • github.com/dmazumdar186

AI Product Manager and hands-on builder — 15 years in data-intensive product. I ship production AI systems end-to-end: GenAI/RAG features, multi-agent workflows, cold-email engines, voice AI and full-stack apps, built with Claude Code and live in weeks, not quarters. A rare pairing of product judgment (discovery, AI PRDs, GDPR/governance, cross-BU alignment) and a builder practice (Python, Cloudflare Workers, Modal, LLM orchestration, eval harnesses) that de-risks delivery.

Impact: \$1M+ pipeline generated • 48,000+ emails/mo at 4 %+ reply • +45 % GenAI adoption • -55 % p95 latency • <30-day median ship • 12+ AI systems shipped

PROFESSIONAL EXPERIENCE

AI Product Manager

Wiser Solutions, Paris | Nov. 2022 – Present

- Shipped production GenAI capabilities (triage, recommendation, RAG support, smart alerts) achieving **+45 % feature adoption and -55 % p95 latency**, by orchestrating OpenAI/Claude multi-agent systems (MCP, A2A) with eval thresholds and prompt-cache-aware cost control
- Drove **+40 % BU adoption and +20 % CSAT** on the global GenAI rollout by authoring AI PRDs (API/data contracts, GDPR/privacy-by-design), running technical feasibility reviews, and aligning **5 cross-BU teams** on a shared roadmap with go/no-go gates and rollback plans
- Lifted delivery precision **+25–30 %** and cut sprint ambiguity **~25 %** by standing up usage/quality/drift dashboards and DoR/DoD + pre-prod QA

AI Product Engineer (Freelance)

High-velocity outbound consultancy (confidential), remote | Dec. 2025 – Mar. 2026

- Designed and shipped an autonomous cold-email engine at scale — **48,000+ emails/month at 4 %+ reply rate**, generating **\$1M+ qualified pipeline** and replacing **~\$200K/yr of SDR headcount**, with Claude-powered personalization, reply classification and contextual auto-reply (3-hour hot-lead SLA via Telegram)
- Architected a serverless event-driven stack (Cloudflare Workers + KV cron, Instantly webhooks, GoHighLevel CRM, scheduled Modal jobs) with idempotency keys, 60-day dedup TTL and LLM guardrails
- Productized monitoring & governance: operator dashboard, **/api/health** endpoint, **--dry-run** on every paid path, synthetic canary → silent drift caught before any API-credit impact

Data Product Manager

InfoTnT, Paris | Jun. 2021 – Nov. 2022

- Led product discovery for a data-driven recommendation engine (quant/qual, JTBD) → **-35 % iteration cycles, +25 % post-launch fit**
- Translated business needs into backend capabilities, epics and data contracts with success metrics and integration constraints
- Standardized prompts, templates and versioning procedures for AI components → predictable behaviors and reliable releases; coordinated releases via PRDs/RFCs and functional QA

Senior Data Product Owner

Pitney Bowes Inc, Pune | Apr. – Sep. 2019

- Reduced time-to-market by **-20 %** via cross-squad dependency mapping, tightened release cadence, pre-prod checklists and structured RAID reviews

Senior Data Product Owner

Evolent International, Pune | Jun. 2018 – Feb. 2019

- Improved platform scalability and performance (**+30 %**) by introducing SLA/SLO to arbitrate run/change priorities; embedded governance and QA checkpoints into the delivery flow

Senior Product Owner

Avaya India Pvt Ltd, Pune | Jul. 2015 – Mar. 2018

- Accelerated delivery velocity (+30 %) through stronger Scrum discipline and OKR ↔ roadmap alignment; reduced requirement instability (–25 %) to improve release readiness

2010 – 2015 : Software Engineer (Tata Consultancy Services, Bengaluru | Nov. 2010 – Mar. 2013) → QA Engineer / Release Coordinator (IDrive India, Bengaluru | Apr. 2013 – Jul. 2015) — technical foundations in distributed systems and release management

SKILLS

AI & GenAI :	LLM, RAG, Agentic AI, Multi-Agent Systems, OpenAI Assistants, Claude, MCP, A2A, supervised / unsupervised ML, AI evaluation frameworks (precision, coherence, robustness), non-determinism handling, prompt caching, LLM guardrails, self-iterating audit loops (anneal), model routing (OpenRouter), Cloudflare Workers, Modal cron, Firecrawl
Product Mgmt :	Vision & Roadmap, Discovery (quant/qual, JTBD), PRD, API/data contracts, Backlog, KPI/OKR, A/B testing, DoR/DoD, Enablement & GTM
Governance :	GDPR / privacy-by-design, access control, audit trails, tech sovereignty, AI ethics
Collaboration :	Cross-functional (Data/AI, Dev, CSM, Sales, MLOps, Security), workshop facilitation, stakeholder management
Tools :	Jira, Confluence, Google AntiGravity, N8N, Make, Figma, Miro, Mixpanel, SQL, GenAI APIs (OpenAI, Google, Anthropic)

EDUCATION

MSc International Strategic Business

Toulouse Business School, Paris | 2019–2021

BE Information Technology Engineering

CMR Institute of Technology, Bengaluru | 2006–2010

LANGUAGES

English : Bilingual (C2) • French : Bilingual (C2) • Hindi : Native • Bengali : Native

SELECTED BUILDS

- **CV Optimizer** (live SaaS) — multilingual AI app that scores a CV against a job description and generates an ATS-optimized CV & cover letter; Streamlit + Gemini, deployed on Cloudflare Pages/Workers
- **Two-sided marketplace** (with a partner agency, under NDA) — Stripe Connect payments platform processing **\$85K in the first 3 months**; React + Node + Postgres
- **anneal** (Python CLI) — self-iterating code-audit loop over git diff (classic + adversarial Red-vs-Blue), dual-adapter LLM (Anthropic + OpenRouter), terminates on 2 consecutive zero-bug rounds
- **YouTube Video Analyzer** — frame-by-frame video breakdown: PySceneDetect + perceptual dedup + 3×3 tiling (**–85 % vision tokens**) + multi-model routing; 73 tests, 8 clean audit rounds
- **job_search_v2** — daily multi-source job aggregator (3 live sources), typed Pydantic v2 pipeline with persistent SQLite dedup and a Gemini A/B/C ranker; runs in CI at **€0/month**