Trusted Segments

Describe your audience. Get a segment you can trust, without a single member record leaving the platform.

The administrator types one sentence. The platform returns the filters it used, inside the existing segmenter, and a **reliability split** on every member: current, to confirm, to reactivate. The model only translates the sentence; the data stays home. Proof on their own file: the same sentence returns **2** (naive filter), **26** (normalised), **49** (full logic) members, of whom only **6** are current or to-confirm and worth e-mailing today.

For the administrator

A list in minutes instead of an hour, and a list that is true: the people who moved are flagged before the send, not after.

For the client's leadership

Dues campaigns reach people who still exist; board numbers carry their reliability; one hour per list becomes minutes; the DPO gets a one-sentence answer to instruct on.

For AlumnForce

The first self-funding step toward the 2027 "platform that acts on its data": a renewal argument for 350 tenants, a sales answer to "do you do AI?", metered per tenant.

The thesis, and the difference

The request says speed; eight of twelve verbatims say trust. To my knowledge, competitors sell the chatbot Antoine asked for: a list. We sell a list **and how far to trust it**, with the one sentence Philippe's DPO accepts. The copyable part is the sentence box. The moat is the normalised, reliability-scored member base per tenant.

15 slides + Annex A (AI usage) · B (selection logic, run on the CSV) · C (technical and compliance sheet) · D (operations and feedback loop)

The segmenter is better than its reputation, and blind to reliability

WHAT I SAW (USERS → LISTS ON THE DEMO INSTANCE)	WHAT IT MEANS FOR THE ASSISTANT
~45 criteria in 9 groups (Account, Dues, Education, Professional situation, Geographic situation, Mentoring…), added one by one through a searchable picker	A fixed vocabulary the model can be constrained to. The brief's "around sixty criteria" is the ceiling, not the floor.
Combination is flat: values joined by **OR** inside a property, properties stacked with **AND**. No nested groups, no relative dates ("last 30 days").	The assistant must say when a sentence cannot be expressed ("finance *or* consulting *unless* Paris") instead of pretending. Saved segments built on "3 years ago" silently age: in Q1 the saved segment stores the relative duration and re-resolves whenever it is opened or sent; a native relative-date operator is a Q2 engine change.
City (workplace) supports **"within 200 km of Lyon"**; Experience (date) supports "is after 1 January 2024"; Sector and the NAF industry code exist on the current job.	"In London" maps to a geocoded radius, not to eight spellings. "For three years" maps to an experience start date. Both already exist.
No criterion on last login or last profile update. Not shown on the member profile either (only a "manually added on" date).	The data that decides whether a list is true is in the database but unreachable in the tool. Exposing it means indexing two existing dates and reusing the existing date filter (slide 6).
Live count only as the table footer ("1-25 de 53"). Lists auto-save as "List - 12/10/2026"; 31 exist, most unnamed. Actions: Export only; newsletters, SMS and dues calls pick lists from their own module.	Yessica's abandonment (V7) is visible in the list names. The assistant should name and describe the segment it creates.
A Profile Update Campaign module already exists; a support chat widget exists; no AI in the segmenter.	The "please update your profile" step reuses an existing module. Less to build.

Also: 12 verbatims coded by need; CSV profiled row by row; Annex B executed in SQLite. Method and verification in Annex A.

Eight of twelve verbatims are not about speed

NEED	WHO SAYS IT	WHAT THEY ACTUALLY WANT
Speed	Sophie V1, Antoine V2, Yessica V7	Say the target once, get the list. Stop falling back to Excel.
Truth	Nathalie V3, Julie V5, Inès V9	A list that does not contain people who left London in 2019. "Nine tickets out of ten are the profile, not the filter."
Proof	Karim V4, Isabelle V10, Catherine V12	Numbers that survive a board meeting; euros justified by time saved and dues collected.
Data	Thomas V8, Philippe V11	"Where does my data go?" answered in one sentence, precisely.
Not for me	Marc V6	400 members, knows everyone. Wants the Excel columns in the right order.

The trap in the request

"Find me the graduates working in finance in London for more than three years." A sentence-to-filters chatbot answers Antoine and fails Nathalie, Julie, Inès, Karim, Isabelle, Philippe and Catherine. It makes the *wrong list faster*.

What is really at stake

- **Administrator**: an hour per list today, and after the send, the embarrassment ("this is nonsense", two unsubscribes).
- **Client leadership**: the network's value is unproven; "roughly" does not pass a board; the DPO can veto the whole feature.
- **AlumnForce**: the segmenter is shown at onboarding and abandoned by month six (V7); one support ticket in five is about lists (V9); large accounts ask "do you do AI?" (V8); the 2027 ambition depends on data the platform can trust.

Coding is mine alone (one coder): speed 3, truth 3, proof 3, data 2, not for me 1; "eight" = truth + proof + data.

The data agrees with Julie: the profile is the problem (n = 500 after deduplication unless stated)

FINDING IN PROFILS-EXERCICE-PM.CSV	VALUE
Rows vs the 500 announced	**525** — 25 apparent duplicate persons (M0501–M0525: surname in capitals, truncated first name, gmail address, same job and start date), 500 after deduplication
Ways of spelling "London"	**8** (london, LONDON, London, Londres, Londres (Royaume-Uni), Greater London, London, UK, London - Canary Wharf)
Ways of spelling the UK	**3**, plus 41 members with no country at all
Declared sector	missing on **174** members; **70 of 127** people at banks declare "Services", "Other" or nothing
"Master in Finance" graduates	47, of which only **19** work in finance (declared sector, employer or job title) — the diploma is not the sector
Profile not updated for 2+ years	**68 %**
No login for 12+ months	**77 %** (89 % if the 60 future-dated logins are treated as unknown)
Last-login dates in the future	**60** members (up to Oct 2027): an export or integrity defect, routed to a data-integrity ticket owned by the data owner, and ignored so it can never count as "recent"

The insight that shapes the product

In this file, profile updates cluster just after a job start (median 112 days, interquartile 50–165; 2 % precede it) and then stop: a correlation, read as "people update when they move, then go silent". So **"in post for 3+ years" and "profile untouched for 3+ years" are the same signal.** All 43 "to reactivate" matches in Annex B have a profile older than 3.9 years and no login in two years.

A naive assistant answering Nathalie's request returns precisely the people most likely to have moved. That is her 340-person "London expats" email, reproduced on this file.

Same sentence, three answers

Naive filter (city value "Londres", sector value "Finance", job started 3+ years ago): **2** members. Normalised city and sector: **26**. Full logic with dedup and three finance sources: **49**. Current or to-confirm, the only ones to e-mail today: **6**, of which 1 current. Naive count on the file as delivered (3 without the tenure clause); the rest on the deduplicated base. (Annex B)

All figures computed from the file with the reference date 2026-10-12.

A reliability split on every segment, not a chatbot

Three deliberate choices

- **Deterministic first.** Normalisation, duplicate detection and the reliability split run in MariaDB/PHP with no model call. They create value for every tenant, including Marc's 400-member club and Philippe's university that will refuse any transfer.
- **The model does one thing.** It turns the admin's sentence into a structured filter set, constrained to the segmenter's existing criteria. It never sees a member.
- **The human keeps the hand.** The filters land pre-filled in the segmenter the admin already knows. They edit, see the split, then save or send. Nothing is sent automatically.

Vocabulary, chosen on purpose

Not "chat", not "grade", not "stale": those read as a reproach to the client's data stewardship. Two names only: the tab is **Trusted Segments**, the chips are **Reliability**: Current · To confirm · To reactivate . Admin vocabulary never reaches members: on their own profile they see "Profile last updated 4 years ago · Confirm / Correct", one button, no label.

What the admin sees (the demo must end on this)

Types: *"alumni from the 2015 to 2018 classes who work in finance in London"*

Gets, in under 5 seconds: the segmenter opened with **Class year: between 2015 and 2018 · Sector (current job): Finance OR Banking OR Finance/Insurance · City (workplace): within 30 km of London**, the part it could not map shown inline under the box ("Add this criterion by hand"), never a blank page, and three clickable chips above the table (illustrative numbers): **72 members** — 31 current · 14 to confirm · 27 to reactivate

One click: send to current + to-confirm; "Reactivate 27 members" opens a pre-filled Profile Update Campaign (the existing module; fixed, tenant-editable text, no model-written copy) and returns here. The segment is saved as "Finance · London · classes 2015-2018" with the sentence as description and a line "Computed on 15/10/2026 · relative dates recalculated when opened or sent", so Karim can quote it to his board next month.

Why not the chatbot Antoine asked for

A free chat over profiles needs member data in the prompt (DPO veto), has no evaluable output, and sets a "ChatGPT" expectation 3 developers cannot meet in a quarter. A constrained assistant with visible filters is evaluable, explainable and reversible.

Every output is one of three things the admin can act on: send, confirm, reactivate.

Four deliverables, a 26 dev-week core sized for shared developers

DELIVERABLE	WHAT SHIPS	DEV-WKS	WHY Q1
1. Reliability layer deterministic · core	Nightly tag: reliability status per member (profile age, last login; future dates ignored) in one MariaDB column and, in OpenSearch, two date fields plus one `reliability_status` keyword; **two new criteria** "Last profile update" and "Last login" through the *existing* date filter, split counts as one terms aggregation: no new criterion type, no cluster refactor; status visible to the member, contest via support ticket in Q1. Reference tables with per-tenant overrides, `*_norm` columns, idempotent backfill.	9	The product itself, on native OpenSearch date filters: a mapping update and a per-tenant reindex, not an engine migration.
2. Sentence → filters one model call · core	Gateway integration; filter contract (Annex C); validator that maps values locally and returns the unmapped part; "3 years" stored as a relative duration plus its resolution date; the shared list loader used by the segmenter and the campaign modules re-resolves it on open and at send time; a manual date edit freezes it and says "fixed date" (no new engine operator in Q1); prompt cache keyed on sentence, tenant, taxonomy, schema, prompt version, role; Vue pre-fill.	9	The "five seconds" value: one schema, zero member data.
3a. Actions core	"Send to current + to-confirm"; "Reactivate" = pre-filled Profile Update Campaign (existing module; fixed, tenant-editable text in two variants compared on update rate); one reactivation per member per quarter; suppression after "not me" or opt-out.	2	The priced lever and the week-17 decision depend on it.
3b. Hygiene stretch / Q2	Duplicate **detection report** (candidate pairs, score, exportable); stable export column order (Marc, V6); relative-date operator as a proper engine change. No merge UI.	3 + 4	Not needed to decide the quarter; the operator is an engine migration and gets its own Q2 rehearsal.
4. Flags, metering, evals core	Two flags, `reliability` and `nl_targeting`, per instance and in bulk, off by default; per-call log (tenant, pseudonymised admin, prompt hash, redacted sentence, tokens, cost, model, region); kill switch; record/replay evals in CI plus nightly replay against the live endpoint; campaign outcomes by segment × status; per-tenant config table (thresholds, radius, allow-list, prompt version) with owner and change log.	6	Cost tracking and stop criteria need it from day one.

Capacity check

3–4 developers **shared** with other tracks, so I plan on roughly half their time: **24–28 effective dev-weeks**, confirmed in week 1. Core = 1, 2, 3a, 4 = 26. 3b and the relative-date operator are Q2 unless capacity appears; below 24, the box pre-fills from a saved-list import instead of live Vue state. The reliability tag and metering never move.

Flag-off promise, kept

Both flags off: nothing changes in the segmenter; the status column exists but no criterion, count or box is shown. `*_norm` columns never overwrite source fields.

No engine criterion type, no cluster refactor, no new operator in Q1: native date filters and a nightly tag. Dev-weeks are mine to be challenged; the Lead Developer's number wins.

Not in Q1, and why

EXCLUDED	REASON	WHEN
Free-text, multi-turn chat over the member base	Requires member data in the prompt; not evaluable; sets a "ChatGPT" expectation (V2) we cannot meet and the DPO (V11) will refuse.	Not planned as such
Duplicate **merge** UI, full data-health panel, relative-date engine operator	Merge touches every module and is irreversible; needs candidate table, reviewer, undo window. The relative-date operator is an engine migration that deserves its own rehearsal. A reliability score carries the pilot without them.	Q2
Retrieval (RAG) or semantic search over profiles in OpenSearch	Only useful once sectors, cities and employers are normalised; otherwise it embeds the noise. Also a transfer of member text to a model.	Q2 at the earliest
Automatic correction of profiles	The platform asks ("still at Lehman Brothers?"), the member confirms. Silent edits destroy the trust we are building.	Never as auto-edit
Enrichment from external sources (LinkedIn, company registries)	Legal basis unclear for most tenants; cost per profile; not needed to prove the increment.	Decision after pilot, per tenant
"Three questions for the board" analytics assistant (V10)	The most valuable ask, but it needs the reliability layer first or it produces confident wrong numbers.	Q2 headline
Automatic sending, SMS, dues calls triggered by the assistant	Human stays between the list and the send. Existing actions on a saved segment are reused unchanged.	Unchanged behaviour
Fine-tuning, own model hosting, multilingual beyond FR/EN	No evidence it is needed; prompt + schema + evals is enough and keeps the provider swappable.	Revisit with eval data

Each exclusion has an owner-level reason, not a capacity excuse.

What goes through the model, and what never does

Human · Deterministic (PHP / MariaDB / OpenSearch) · Language model, via the internal gateway

1 · ADMIN
Types a sentence in the segmenter. Session carries tenant, role and visible criteria.

→

2 · PRE-FLIGHT
Flag and quota check; length cap; scrub e-mails, phones and exact full-name matches (first + last token) from this tenant's member table; replace any other capitalised two-token sequence by a placeholder the admin sees, with place and sector words allow-listed; log every hit, never block; false-flag rate measured on the eval set; cache lookup.

→

3 · MODEL CALL
Input: scrubbed sentence + the tenant's full criteria schema (names and operators only, no data) + **CLOSED TAXONOMIES ONLY** (sector list, country list, diploma types). Output: strict JSON filter set with place text and relative durations, or "cannot map" with the unmatched part.
NO MEMBER ROW, NO EMPLOYER LIST, EVER.

→

4 · VALIDATOR
Rejects unknown criteria; geocodes place text and resolves durations locally; maps free values to reference tables; sanity bounds (radius cap, operator/value types per criterion, result-count vs tenant-size alert); one repair retry then "cannot map"; hard timeout → degraded; logs tenant, admin, tokens, cost.

→

5 · QUERY + SPLIT
Existing segmenter engine runs the filters on MariaDB/OpenSearch; each result carries its reliability status.

→

6 · ADMIN DECIDES
Sees filters pre-filled next to their own sentence (a diff view), the split, the unmapped part inline. Edits, saves, sends, or reactivates.

Rights

The assistant runs inside the admin's session. The model sees the tenant's criteria names, never data, so the cached prefix is shared per tenant; the validator then strips any criterion the role may not use (unit-tested, not left to the model), and results are scoped by the existing query layer, and the cache key includes tenant, taxonomy hash, role and schema hash so no filter set leaks across roles or instances.

Where the human is

Before (writes the sentence), in the middle (sees and edits the filters), after (chooses the action). The model never triggers a send, a merge or a profile edit. Prompt injection in the sentence can produce a rejected filter set or a wrong one (for instance widened to everyone), which the admin sees with its count before any send; widening attacks are in the negative eval set.

Degraded mode

Gateway down, quota reached, flag off, timeout: the segmenter works exactly as today. Partial "cannot map": mapped filters apply, the rest is shown inline, never a blank page. Global kill switch at the gateway, owned by the Lead Developer. The feature is additive.

The model is a translator with a fixed vocabulary, not a search engine over people. Filter contract and failure modes in Annex C.

What the DPO can instruct on, and what a use costs

For Thomas (V8) and Philippe (V11), one sentence

"Only the administrator's question, scrubbed of identifiers and full names, and the names of the filters leave the platform, to a model hosted in the EU under a no-retention, no-training contract. No member profile is ever sent; residual hits in the question are logged and reviewed."

The file behind the sentence (Annex C has the questionnaire rows)

- **Roles**: tenant controller, AlumnForce processor, model provider a sub-processor with 30-day notice; DPA with zero retention, no training; EU hosting *and* jurisdiction stated.
- **DPIA**: AlumnForce supplies a screening memo and a balancing-test template; each tenant, as controller, records its own basis and decision; the status can be disabled per tenant.
- **Reliability status is derived personal data.** Computed per tenant only after the flag is on, never for all 350 by default. Likely basis: the tenant's legitimate interest in accurate records, the tenant's call. Shown to the member on their profile, contestable in one click (a support ticket to the tenant admin in Q1, answered within 10 business days; a contest sets the keyword to "confirmed" and freezes it for 12 months); template notice clause. Not an export field, though a status-filtered export reveals it, and the notice says so. Reactivation honours opt-outs and the tenant's basis for contacting silent members.
- **Logs**: pseudonymised admin, prompt hash and redacted sentence, 12 months, tenant-scoped export, deleted on offboarding; saved sentences follow the same policy.

Cost of one use (confirmed on the gateway in week 1)

~1,500–2,500 input tokens (schema and taxonomies, in a cached prefix), ~150–300 output: a fraction of a euro cent per request on a mid-tier model, under one euro a month for a heavy tenant. Per-admin rate limit, per-tenant quota with soft block and CS alert. The cost that matters is engineering (slide 13).

Provider position

Contract is "schema in, JSON out", so the provider is swappable in days. If today's gateway points to a US provider, the sentence box launches only on an EU-hosted endpoint; the reliability layer does not wait for that decision.

Security questionnaire, ready

One row per item in Annex C (TLS, gateway auth, certifications, residency, pentest scope, incident SLA), so Sales never answers "it's secure". No-transfer tenants keep everything except the sentence box.

Minimised, justified, traceable, refusable: the brief's four words, each with a mechanism.

Quality, adoption, value, and one stop criterion per layer (every threshold is an initial guess, reset on the week-4 baseline)

Quality, offline (before any tenant sees it)

- Golden set of **100 real sentences** from tickets and the verbatims, labelled by two people independently; inter-annotator agreement reported. Negative set: sentences that must return "cannot map" (discriminatory, impossible, injection attempts).
- Metrics: filter-level precision and recall, exact-match rate, "cannot map" accuracy. Gate: exact match ≥ 85 % with its 95 % interval stated (on 100 items, roughly ±7 points), zero unknown criteria or values after validation; role stripping is covered by unit tests, not evals. Re-run on every prompt, schema or model change, and **replayed nightly against the live endpoint** with an alert on a 5-point drop, because providers update silently; provider model version logged per call.

Quality, online

- **Edit rate** (filters changed before saving, per segment), **rejection rate** (typed, nothing saved, per sentence), **cannot-map share** (per sentence): all per tenant, aggregate and pseudonymised, disclosed to admins at flag-on. Weekly sample of 20 sentences, never attributed to a person. Edited sentences become eval candidates.

Adoption (per tenant, visible to CS, not only to me)

- Share of new segments created via the sentence box; weekly active admins; return use in weeks 2 and 4.
- Excel export share of list-building (should fall); list-related tickets per 100 campaigns (target −30 % on pilot tenants, tagged from week 1).

Value

- Time from first click to saved segment (baseline weeks 1–4).
- "Not me / I moved" replies and unsubscribes per campaign (V3).
- Share of "current" profiles per tenant month over month; profile updates triggered by reactivation campaigns.
- Response rate on dues campaigns sent to current + to-confirm vs historical.

Stop criteria, one per layer, two dates

- **Sentence box, week 12** (3 pilot weeks, leading indicators): stop if fewer than **30 %** of new segments on the 7 non-expert tenants use it, **or** edit rate stays above **50 %**, **or** the eval gate is not met. Keep the reliability layer.
- **Reactivation loop, week 17** (30 days after the first campaign): stop it if fewer than **10 %** of "to reactivate" profiles move to "current". Keep the status and criteria, which cost nothing to run.
- Expert tenants are reported separately so three power users cannot carry the number.
- **Owners**: quota and anomaly alerts → CS on-call, one business day; eval gate or nightly replay failure → PM blocks release, Lead Developer notified; integrity tickets → data owner, weekly.

Baselines are measured on the current segmenter before launch, so the pilot compares against a number, not a feeling (V7).

Three risks, and my position on each

1 · The data is too dirty for any assistant to look good

68 % of profiles untouched for two years; sectors missing on a third; eight spellings of one city. An assistant on top of this returns wrong lists with confidence, and the brand damage lands on AlumnForce, not on the model.

Position: the reliability layer *is* the first increment, not a prerequisite someone else owns. The split makes the dirt visible instead of hiding it, and reactivation starts cleaning it from day one. I accept shipping the sentence box later than a pure-AI roadmap would.

2 · Personal data and the DPO veto

Philippe (V11) will block any transfer of student data to a US model; several tenants will refuse categorically; the reliability status itself is profiling if nobody says so. One ambiguous answer in a sales call and the feature is dead for that account.

Position: design so the answer is short and true: scrubbed sentence only, closed taxonomies, EU endpoint under a no-retention contract, status visible and contestable by the member, per-tenant flag and audit export. Tenants who still say no get everything except the sentence box. I will not trade this for a richer prompt.

3 · The expectation gap

Antoine (V2) expects ChatGPT. The product is a translator with a fixed vocabulary. If the first demo feels like a dumber chatbot, pilots disengage and Yessica's story repeats (V7).

Position: never call it a chat; always show the filters and the unmapped part; make the first win the split, which no chatbot gives them; end every demo on "6 current or to-confirm you can e-mail". Onboard the three expert tenants first so the pattern is set by people who know what a filter is.

Watched but not top three: cost runaway (metered, capped, kill switch); provider lock-in (schema-in/JSON-out contract); scope slip (deliverable 3b and the relative-date operator are already Q2; the Vue pre-fill moves next); small tenants finding it irrelevant (V6: flags off, they still get the normalisation backfill and a reliability score at renewal; the duplicate report follows in Q2).

Each position is a trade I am willing to defend, not a mitigation list.

EXECUTION

Twelve weeks, four weeks of measurement, then waves

WEEKS	WORK	OUTPUT
1–2	Instrument the current segmenter (time-to-list, exports, ticket tags). Golden set with CS. Thresholds validated first by hand: CS grades 100 members on 2 tenants; the 2 y / 4 y rule must agree ≥ 80 % before the nightly job ships. DPIA screening memo, DPA and sub-processor check, gateway region and FTE confirmed.	Baselines; golden set v1; legal file opened
3–7	Deliverable 1: nightly status tag, two date fields plus a status keyword in the OpenSearch mapping, criteria through the existing date filter, per-tenant reindex, reference tables, chunked backfill (80k instance in one night, measured on staging week 5). Deliverable 4 in parallel.	Reliability split live behind the flag, week 7
6–9	Deliverable 2: contract, gateway, validator, caching, Vue pre-fill; eval gate before week 10. Deliverable 3a (actions); 3b and the relative-date operator only if capacity appears, otherwise Q2.	Sentence box behind the flag, week 9
8–12	Pilot in two steps: reliability layer from week 8, sentence box from week 10. Ten tenants: 3 expert, 5 mid, 2 small (one refusing transfer, to prove the no-model path), including 3 top-decile ARR accounts and every renewal flagged at risk in the next 9 months; high list-ticket volume, DPO already consulted. Pilot agreement: weekly 30-minute call, baseline access, a quotable testimonial, a named case study.	Pilot dashboard, visible to CS
12 · 17	Week 12: go / no-go on leading indicators (eval gate, adoption and edit rate, ticket trend). Week 17: reactivation outcome 30 days after the first campaign; pricing test. Q2 brief follows: duplicate report, relative-date operator, merge UI.	Decision memo, then a confirmation memo
Q2+	Waves of ~50 tenants a month, at-risk renewals first; a wave is backfill plus flag, zero new development. Backfill off-peak, CS notified. Flag-off tenants get the normalisation backfill and, if 3b ships, the duplicate report and export column order (V6); the status is computed only once a tenant has recorded its DPIA decision and switched the flag. Enablement in week 10: one-pager, FAQ, demo script, objection sheet.	350 tenants in ~7 months

Week 1, with or without inputs

I will ask for ticket data, the criteria model, the gateway's provider and region, and the tenants who refused AI. Without them I assume: a 200-ticket hand-tagged sample, role-level permissions, a US-owned EU region, ten refusing tenants; and I keep going. Developer split is in the speaker notes.

Two decision points because the two layers prove themselves at different speeds; the quarter still ends with a decision.

MONEY

What pays for the quarter, and what we charge

Three levers, one primary

- **Retention (primary)**: 350 installed tenants, segmenter abandoned by month six (V7). A reliability score at renewal is a CS argument I have not seen a competitor make.
- **New logo**: "do you do AI?" (V8) answered with the sentence on slide 9 and Annex C.
- **Upsell**: the reactivation loop, the only thing tied to dues collected (V12).

Payback, as formulas with the baselines measured in weeks 1–4

- **Cost**: 26–28 developer-weeks at your loaded rate; inference rounds to zero.
- **Tickets**: list tickets/month × avoidable share (target 30 %) × CS cost per ticket.
- **Admin time**: segments/month × minutes saved × admin cost; the client's gain, so it feeds renewal, not margin.
- **Dues**: profiles moved to "current" × response uplift × average dues.
- **Breakeven I accept**: engineering recovered within four quarters from ticket savings and retained ARR across the pilot and the first 50-tenant wave, before any upsell.

Packaging hypothesis, tested in the pilot

- **Included for all**: status, criteria, split counts, score at renewal; duplicate report when it ships in Q2.
- **Included with fair use**: the sentence box, 300 requests/month, soft block and CS alert beyond. Catherine will not pay for a text box (V12).
- **Priced add-on**: the reactivation loop, tested on the 3 expert tenants as a flat monthly fee per member tier, a fraction of the dues it demonstrably recovers, reverting to included if reactivation stays under 10 %. Proof on the client's own file: 2 → 26 → 49 → 6 current or to-confirm.

Renewal pack, generated per tenant each quarter

Score and its delta, hours saved, "not me" replies before and after, dues response of current + to-confirm sends vs historical, and an attribution line (reactivations sent → profiles updated → dues recovered) so the uplift is credited to the platform. Isabelle's "prove the value of the network" (V10) as a by-product.

Competitive line

Hivebrite, Graduway and similar market AI search and assistants; to my knowledge none scores a segment's reliability or states in one sentence what leaves the platform (to verify in week 1). Press hypothesis: "an AI that is never sent a profile, hosted in Europe". The sentence box is copyable in a month; a normalised, reliability-scored, consented base per tenant is not.

Numbers in the formulas are targets to measure, not forecasts; I will not quote a euro figure I have not seen.

Each quarter funds the next, or stops

Q2 — Answers for leadership

"How many alumni placed, where, in which sectors" (V10) as three questions over the reliable base, on the existing Statistics module's data: every number carries its reliability. Same translator, aggregate output, still no member data in the prompt. Funded as the renewal argument for client directions. Also: merge UI with undo window, data-health panel.

Q3 — Automations that keep the base alive

n8n loops: reactivate on job-change signals, re-evaluate status after each login, propose merges. Funded by the priced reactivation loop if the pilot proved it.

2027 — Agents on segments

An agent proposes a dues campaign to the right 180 people, with the list, the split and the expected response rate, and waits for a human yes. Possible only because the three previous steps made the data trustworthy and the actions auditable.

Decision I am asking for

Fund one quarter, plus four weeks of measurement, for **Trusted Segments**: reliability layer first, one constrained model call, human in the segmenter, metered per tenant, two flags, one stop criterion per layer, decided at week 12 and week 17.

What you get at week 12

- Ten tenants with a reliability-scored, normalised base and a reliability score at renewal.
- A sentence box that pre-fills the segmenter, with an eval gate and a per-tenant cost line.
- A one-sentence data answer and a questionnaire-ready sheet that Sales and CS can use tomorrow.
- Measured numbers: time to list, edit rate, "not me" replies, "current" share, dues response.

What I would not do even if asked

- Put member rows or employer lists in a prompt to make the assistant "smarter".
- Let the model send, merge or edit anything.
- Ship the sentence box before the reliability split exists.

Questions I bring to the exchange are in the speaker notes; the first is which tenants already said no to AI, and why.

CLOSING

Two sentences to remember

"The administrator's problem is not building the list. It is believing it. So the first thing the platform returns is not a list, it is how far to trust it."

"Only the administrator's question, scrubbed of identifiers, and the names of the filters leave the platform. Never a member's profile."

350 instances. 68 % of profiles untouched for two years in the 500-member sample. The asset is a reliable alumni base, not the text box on top of it.

Debanjan Mazumdar · Senior Product Manager Data & AI · AlumnForce case study · October 2026

Thank you. 30 minutes of your questions, then 10 minutes of mine.

How AI was used, and what was checked

What AI did (most of the volume)

- Read the brief and verbatims with me and proposed the coding of needs; drafted every slide from my outline and rewrote them through five review rounds.
- Wrote and ran the pandas and SQL that profile the CSV and produce Annex B; found the 25 apparent duplicates and tuned the matching rule until each matched exactly one original.
- Ran two automated browser sessions on the demo instance under my admin account: the first listed menus and criteria and created then deleted one empty test list; the second was read-only and re-checked the criteria picker, the Profile Update Campaign module and the list count. No screenshot appears here.
- Played a recruiter panel and eleven reviewers (measurement, tooling, honesty and law, rigor, offer, automation, people, failure modes, UX, reliability, strategy) against each draft, then re-ran them until clear.

What was verified, and by whom

- Annex B re-run on the CSV: 49 rows, the 2 / 26 / 49 / 6 cascade and every slide-4 figure, by an agent on my instruction and by two independent reviewers; the record with commands and outputs is attached (VERIFICATION.md).
- The SQL printed on the Annex B page and the attached file were checked to be identical and to return the same 49.
- The demo-instance facts (45 criteria in 9 groups, no last-login criterion, Profile Update Campaign present) by the read-only session.
- Every position, exclusion, threshold and stop criterion is my decision; AI proposed, I chose.

Where I did not trust it

Anything about AlumnForce internals, competitors, provider contracts and legal bases is stated as an assumption or a question, not a fact. Token costs are an order of magnitude. Payback is formulas, not euros. Where a review lens could not verify something, its honest gap is kept rather than papered over.

Discernment, not abstinence: AI for volume and for checking, me for every claim I sign.

ANNEX B · SELECTION LOGIC (ONE PAGE) · THE BRIEF'S REQUEST: "GRADUATES WORKING IN FINANCE IN LONDON FOR MORE THAN THREE YEARS" · FULL SQL AND 49-ROW RESULT ATTACHED (ANNEXB.SQL, ANNEXB_RESULT.CSV)

```
-- Reference date 2026-10-12. Runs as written on SQLite (LIKE only, no REGEXP); for MariaDB replace date(d,'-N
years') by DATE_SUB(d, INTERVAL N YEAR).
-- Identical to the attached annexB.sql (same keyword list); column names are the dataset's own (French).
-- 0. DEDUPLICATE: file has 525 rows, brief says 500. M0501..M0525 are duplicate persons
--    (capitalised/compound surname, truncated first name, gmail, same job + start date).
CREATE TABLE members AS            -- (preceded by DROP TABLE IF EXISTS in the attached file)
WITH p AS (SELECT *, lower(substr(prenom,1,3)) fn,
           lower(replace(replace(replace(nom,'-',''),' ',''),'.','')) ln FROM profils),
duplicates AS (SELECT b.member_id drop_id FROM p a JOIN p b
   ON a.member_id<>b.member_id AND a.fn=b.fn
  AND (instr(a.ln,b.ln)>0 OR instr(b.ln,a.ln)>0)
  AND a.date_debut_poste=b.date_debut_poste                -- same start date: Martin is not Martinez
  AND (a.date_maj_profil>b.date_maj_profil OR (a.date_maj_profil=b.date_maj_profil AND a.member_id<b.member_id)))
SELECT * FROM p WHERE member_id NOT IN (SELECT drop_id FROM duplicates);    -- 500 members
-- 1. NORMALISE (deterministic lookup tables, no model, nothing leaves the platform)
CREATE TABLE members_norm AS SELECT m.*,
 CASE WHEN lower(ville) LIKE '%london%' OR lower(ville) LIKE '%londres%'
       OR lower(ville) LIKE '%canary wharf%' THEN 'London' ELSE ville END city_norm,   -- 8 spellings
 CASE WHEN lower(trim(coalesce(pays,''))) IN ('uk','united kingdom','royaume-uni','england','angleterre') THEN 'GB'
      WHEN pays IS NULL AND lower(ville) LIKE '%lond%' THEN 'GB' ELSE pays END country_norm,
 CASE WHEN secteur_declare IN ('Finance','Banque','Finance / Assurance') THEN 'finance:declared'
      WHEN entreprise IN ('Amundi','AXA IM','Barclays','BNP Paribas','Crédit Agricole CIB',
        'Deutsche Bank','Goldman Sachs','HSBC','J.P. Morgan','Lazard','Morgan Stanley',
        'Natixis','Rothschild & Co','Société Générale','UBS') THEN 'finance:employer'
      WHEN lower(poste) LIKE '%financ%' OR lower(poste) LIKE '%m&a%' OR lower(poste) LIKE '%investment bank%' OR
lower(poste) LIKE '%ib analyst%'
        OR lower(poste) LIKE '%trader%' OR lower(poste) LIKE '%quant%' OR lower(poste) LIKE '%portfolio%'
        OR lower(poste) LIKE '%asset manag%' OR lower(poste) LIKE '%opcvm%' OR lower(poste) LIKE '%private equity%'
        OR lower(poste) LIKE '%equity research%' OR lower(poste) LIKE '%treasury%' OR lower(poste) LIKE '%credit
risk%'
        OR lower(poste) LIKE '%risk manager%' OR lower(poste) LIKE '%contrôleur de gestion%' OR lower(poste) LIKE
'%fp&a%'
        OR lower(poste) LIKE '%banquier%' OR lower(poste) LIKE '%conformité bancaire%' OR lower(poste) LIKE
'%auditeur financier%'
        OR lower(poste) LIKE '%chargé d''affaires%' OR lower(poste) LIKE '%gestionnaire de portefeuille%' OR
lower(poste) LIKE '%directeur financier%' THEN 'finance:title'
      ELSE 'none' END finance_src,                       -- job-title keywords are French because the data is
 -- bounds: A = [0, 2 y) since update OR [0, 1 y) since login; B = [2, 4 y); C = 4 y+ or no date
 CASE WHEN date_maj_profil > date('2026-10-12','-2 years')                                  -- current
       OR (derniere_connexion > date('2026-10-12','-1 year') AND derniere_connexion <= '2026-10-12') THEN 'A'
      WHEN date_maj_profil > date('2026-10-12','-4 years') THEN 'B'                        -- to confirm
      ELSE 'C' END reliability                                                              -- to reactivate
FROM members m;
-- 2. SELECT
SELECT member_id, poste, entreprise, finance_src, reliability FROM members_norm
WHERE city_norm='London' AND country_norm='GB'            -- "in London" (country gate: not London, Ontario)
  AND finance_src<>'none'                                 -- "working in finance"
  AND date_debut_poste <= date('2026-10-12','-3 years')    -- "for more than three years"
  AND CAST(promo AS INTEGER) <= 2026;      -- "graduates": all are (class years 2005-2024); CAST survives a text
import
```

Result

49 members (500 after dedup → 168 in London → 72 in finance → 49 with 3+ years in post). Source: 24 declared sector, 17 employer, 8 job keywords. Reliability: Current 1 · To confirm 5 · To reactivate 43 · 14 paying. Naive filter (city value 'Londres', sector value 'Finance', 3+ years in post) on the file as delivered: **2**. Normalised, no dedup: **26**. "6 to e-mail" = current + to confirm; strictly current: 1.

Assumptions

- Reference date 2026-10-12 (latest non-future date in the file: 2026-10-06).
- "Graduates" = everyone (class years 2005–2024). The "Master in Finance" diploma is not a finance signal: only 19 of its 47 holders (after dedup) work in finance by any of the three sources.
- "Finance" = declared sector, else employer in a finance reference list, else job-title keywords. Keywords are the weakest source: 3 of 8 are doubtful ("Professeur de finance" at BCG and at Renault, "Assistant administratif et financier" at Orange); in the product, the NAF industry code comes first and keyword matches are shown "to confirm".
- "3+ years" = current job started on or before 2023-10-12.
- Dedup rule (same first-name prefix, surname contained, same start date) drops exactly the 25 apparent duplicates; an earlier "or same job" variant also collapsed Martin into Martinez, which is why production uses a scored candidate list with blocking keys, reviewed by the admin.

What the data cannot verify, and what the product does

- **"In London for 3+ years"**: no date for the location, only for the job. Job tenure is the proxy, labelled "3+ years in current role" in the UI.
- **"Still true"**: 43 of 49 profiles are 4+ years old with no login in 2 years. Marked "to reactivate"; a Profile Update Campaign is offered before any send instead of silent inclusion.

ANNEX C · TECHNICAL AND COMPLIANCE SHEET (ONE PAGE) · THE FILTER CONTRACT, FAILURE MODES, AND THE SECURITY-QUESTIONNAIRE ROWS

```
// Filter contract: model output, validated server-side (JSON schema enforced; 1 repair retry, then "cannot
map")
{
  "filters": [
    {"criterion": "class_year",         "op": "between", "values": [2015, 2018]},
    {"criterion": "sector_current_job","op": "in",      "values": ["Finance","Banking","Finance / Insurance"]},
    {"criterion": "city_workplace",    "op": "near",    "place_text": "London", "radius_km": 30},
    {"criterion": "experience_start",  "op": "before",  "relative": {"years": 3}}
  ],
  "unmapped": ["who answered the last survey"]
}
// Rules
// - the tenant's full criteria schema is sent (names/operators only); the validator strips criteria outside
the admin's role
// - tenant-defined criteria (custom fields, event titles) are sent as typed ids ("custom_12: amount") unless
the tenant opts
//   their labels in; standard criteria names are AlumnForce's own vocabulary, not tenant data
// - values for closed taxonomies only (sector, country, diploma type); place_text and employer text are
resolved locally
// - Q1: the segment stores {relative: {years: 3}, resolved_at}; the shared list loader (segmenter, newsletter,
SMS, dues
//   modules) re-resolves on open and at send time on the tenant clock; a manual edit of the date drops
`relative` and the
//   chip reads "fixed date"; a 90-day-stale segment fixture is in the eval set; native operator is a Q2 engine
change
// - cache key = hash(scrubbed sentence) + tenant + taxonomy hash + schema hash + prompt version + model id +
role; TTL 24 h
// - place_text resolved by the segmenter's existing geocoder (the radius operator already uses one); nothing
new leaves
// - scrub placeholders (e.g. {{P1}} for "Goldman Sachs") are re-substituted locally in the validator before
employer/city
//   resolution, so "those who work at X" maps; round-trip cases are in the eval set
// - exports are dated snapshots
// - log row = tenant, pseudonymised admin, prompt hash, redacted sentence, tokens in/out, cost, model, region,
latency
// Reliability status (nightly tag: one MariaDB column; OpenSearch: two date fields + reliability_status
keyword, split = one terms
//   aggregation, criteria via the existing date filter; per-tenant reindex)
//   computed per tenant after the reliability flag is on (post DPIA decision); nightly thereafter
//   thresholds 2 y / 4 y / 1 y are initial (update lag median 112 d in the sample); validated by hand-grading
in weeks 1-2
//   members with no update date and no login -> C; future dates ignored
//   A "current"        profile updated < 2 y  OR  last login < 1 y (future dates ignored)
//   B "to confirm"     profile updated 2–4 y
//   C "to reactivate"  older; exposed as two criteria + counts; shown to the member on their profile,
contestable
// Backfill: chunked, idempotent, *_norm columns only; 80k-member instance measured on staging before any wave
// Failure modes: gateway down / quota / flag off / timeout / schema mismatch -> segmenter unchanged, event
logged;
//   partial cannot-map -> mapped filters applied, fragment shown inline
// Config: per-tenant assistant_config (thresholds, radius cap, allow-list, prompt version, taxonomy hash) with
owner;
//   CHANGELOG row per prompt/schema/taxonomy bump: version, eval delta, incidents; taxonomy bump purges cache
// Eval harness: record/replay fixtures in CI; live run on prompt or model diff; nightly replay vs live
endpoint
//   exact match = canonicalised filter set (sorted criteria, sorted values) equal to label; P/R per
(criterion, op, values) tuple
//   2 annotators, Cohen's kappa >= 0.8; negative-set recall >= 95 %; 100 labelled + 30 negative
//   pinned: model id, temperature 0, prompt version; live runs k=3, flake budget 2 %
```

QUESTIONNAIRE ITEM	ANSWER
Data categories sent to the model	Scrubbed administrator question; criterion and operator names; closed taxonomy values. No member fields, no employer lists.
Roles	Tenant controller; AlumnForce processor; model provider sub-processor (DPA, zero retention, no training; listed with 30-day notice).
Hosting and jurisdiction	EU-hosted endpoint; provider legal entity and jurisdiction named; transfer mechanism stated if US-owned; swappable (schema in, JSON out).
DPIA	Screening in week 1; balancing test documented for the reliability status (legitimate interest of the tenant).
Member transparency	Status visible on own profile, contestable; template clause for the tenant's notice plus in-product notice, adoption logged per tenant; not an export field, though a status-filtered export reveals it (stated in the notice).
Logs	Pseudonymised admin, prompt hash + redacted sentence, 12 months, tenant-scoped export, deleted on offboarding; read access at AlumnForce restricted to named support roles and itself logged.
Access control	Admin session; role-pruned schema; instance-scoped query; cache keyed on role.
Abuse and injection	Rate limit per admin, quota per tenant, anomaly alert; injection yields a rejected or a visibly wrong filter set, never a send; widening attacks in the negative eval set.
Transport and auth	TLS to the gateway; per-instance gateway credentials; kill switch.
Assurance	Provider certifications and residency attestation on file; pentest scope includes the sentence box; incident notification per existing SLA.
Right to refuse	Per-tenant flag; refusing tenants keep the reliability layer; audit export on request.

ANNEX D · OPERATIONS, OWNERSHIP AND THE FEEDBACK LOOP (ONE PAGE) · WHAT MUST BE TRUE FOR THIS TO RUN UNATTENDED FOR 30 DAYS

OPERATION	DESIGN
Nightly status job × 350 tenants	Per-tenant lock and watermark, resume from last tenant, tenant-level success/failure; dead-man alarm if any enabled tenant's status is older than 36 h.
Per-tenant reindex	Alias-swap: build new index, verify counts, flip alias; canary 1 → 3 → 10 tenants; rollback = alias flip.
Mapping change (two date fields + status keyword) in Q1; relative-date operator in Q2	Additive mapping fields that ride the alias-swap reindex, legacy saved lists untouched; the Q2 operator gets an expand/contract migration rehearsed on a staging snapshot.
Backfill	Upsert on member_id with run-id, resumable; idempotency proven by a double run on staging before wave 1.
Taxonomy or reference change	Taxonomy hash bump → cache purge → targeted re-normalise job; recorded in the CHANGELOG.
Gateway credentials	Per-instance, dual-key rotation every 90 days, automated.
Runbook (one page, 6 am)	Symptom → check → action → owner: eval replay drop → flag off for new tenants, PM decides; provider degraded → kill switch (Lead Developer); quota storm → CS on-call; stale status → re-run tenant job.
Members	Contest answered by the tenant admin within 10 business days and freezes the status as "confirmed" for 12 months (relief, never a downgrade); one reactivation per member per quarter; suppression after "not me" or opt-out; CS script for "why am I flagged".

MEASUREMENT	DECISION IT FEEDS (NO DEAD ENDS)
Admin edits a generated filter	Sentence + output + correction become weekly eval candidates; golden set grows from production; edit-rate target tightens each quarter.
"Not me" reply (a contest is the opposite: it freezes "confirmed")	Member flips to to-reactivate with suppression; share of current members saying "not me" recalibrates the 2 y / 4 y thresholds.
"Cannot map" fragment	Logged per tenant; ranked list is the criteria and operator roadmap (relative dates were found by hand this time).
Tenant override of employer → sector	Confirmed by N tenants → promoted to the global table; 350 instances teach one table.
Campaign outcome by segment × status	Stored from Q1; the 2027 agent's "expected response rate" is computed from it, not asserted.
Nightly replay against the live endpoint	5-point drop → release blocked, provider version compared, prompt pinned.
Reactivation text	Two fixed, tenant-editable variants compared on update rate; the priced lever improves or is dropped.

Flywheel metric
Profiles moved from to-reactivate to current per month across enabled tenants, with edit rate falling alongside. If both move, the loop compounds; if neither does, the stop criteria fire.