Trusted Segments

Describe your audience in one sentence. Get a segment you can trust. No member record ever leaves the platform.

You asked for speed. Your users asked for trust.
So the platform returns not just a list, but how far to trust it.

Full case study, annexes and SQL: deck.pdf

THE PROBLEM, REFORMULATED

Eight of twelve verbatims are not about speed

3

want speed
"Saying it takes five seconds, doing it takes an hour."

3

want a true list
"340 London expats. Twenty had moved back years ago. Two unsubscribed."

3

want proof
"'Roughly' does not pass a board."

2

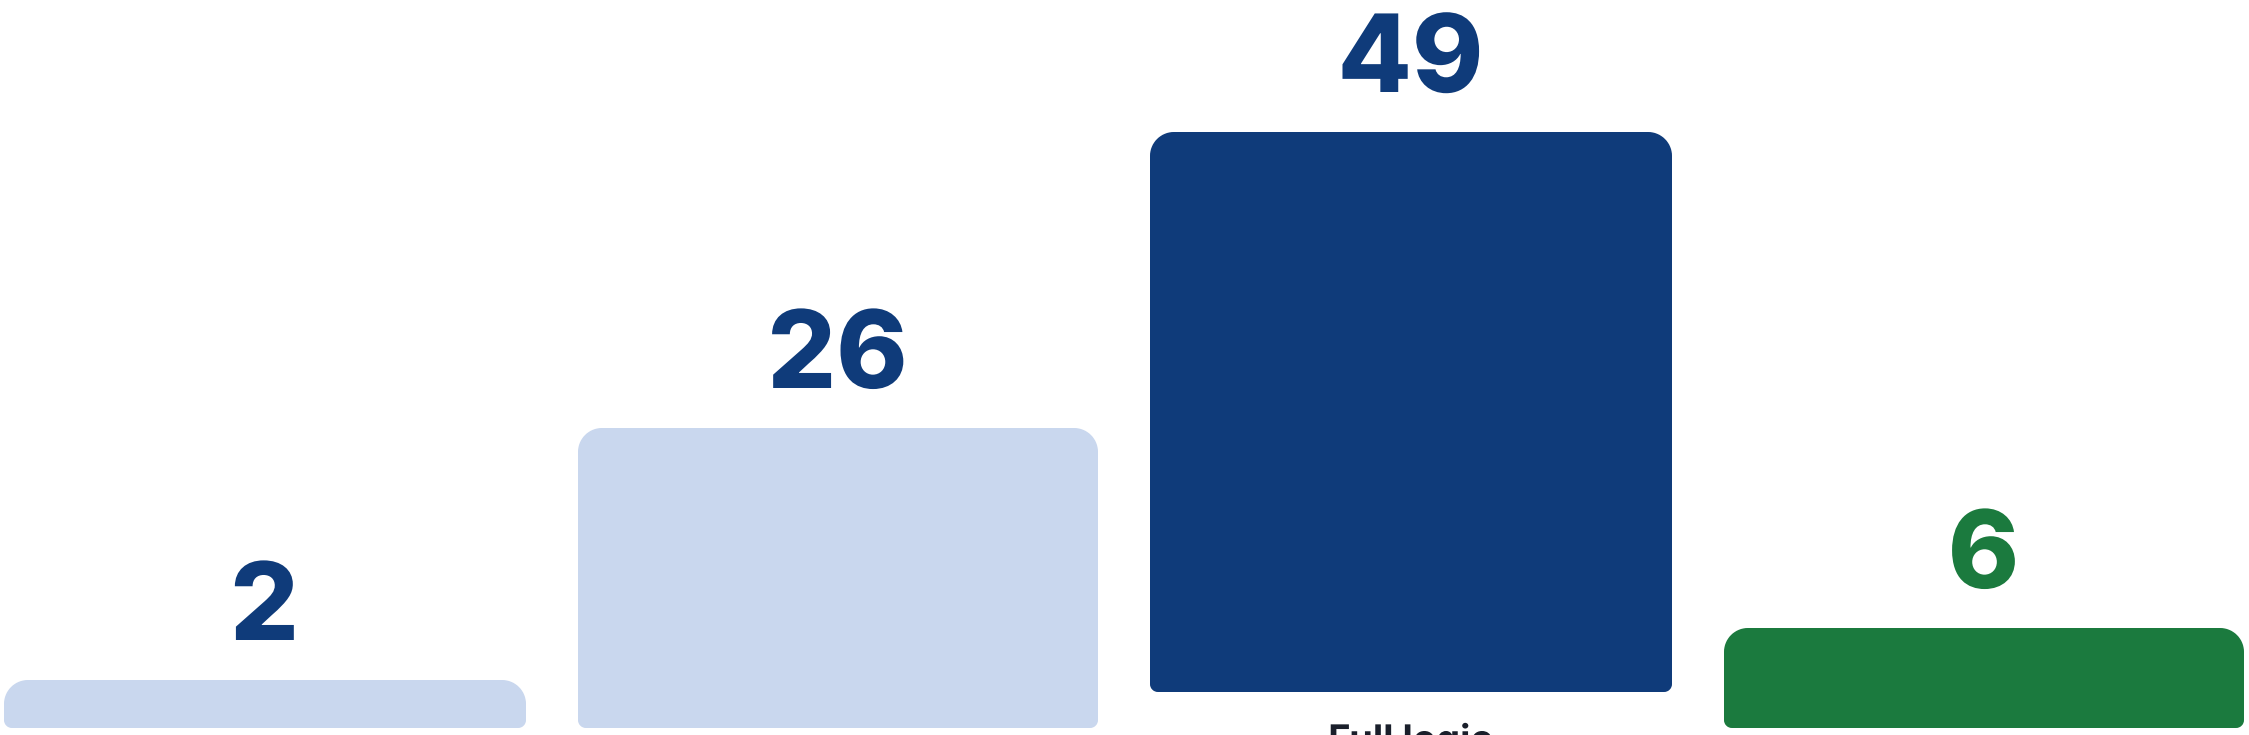
want a data answer
"Tell me exactly what leaves the platform."

A chatbot that turns a sentence into filters serves the first three people. It makes the **wrong list faster** for everyone else.

One member says "not for me": 400 people, knows everyone, wants the Excel columns in order.

Same sentence, four answers

"Graduates working in finance in London for more than three years"



Naive filter

city = "Londres", sector = "Finance", exact text

Normalised

8 spellings of London and 3 sector labels merged

Full logic

duplicates removed, employer and job title used as finance signals

Worth e-mailing today

profile or login recent enough to trust

68 % of profiles untouched for 2+ years · 525 rows for 500 announced members · 8 spellings of London

THE INSIGHT

In this data, "three years in post" and "three years without an update" are the same signal

People update their profile when they start a job (median 112 days later) and then go silent. A naive assistant answering the brief's own example returns exactly the people **most likely to have moved**.

That is the 340-person London e-mail disaster, reproduced on their own file.

So every segment gets a reliability split

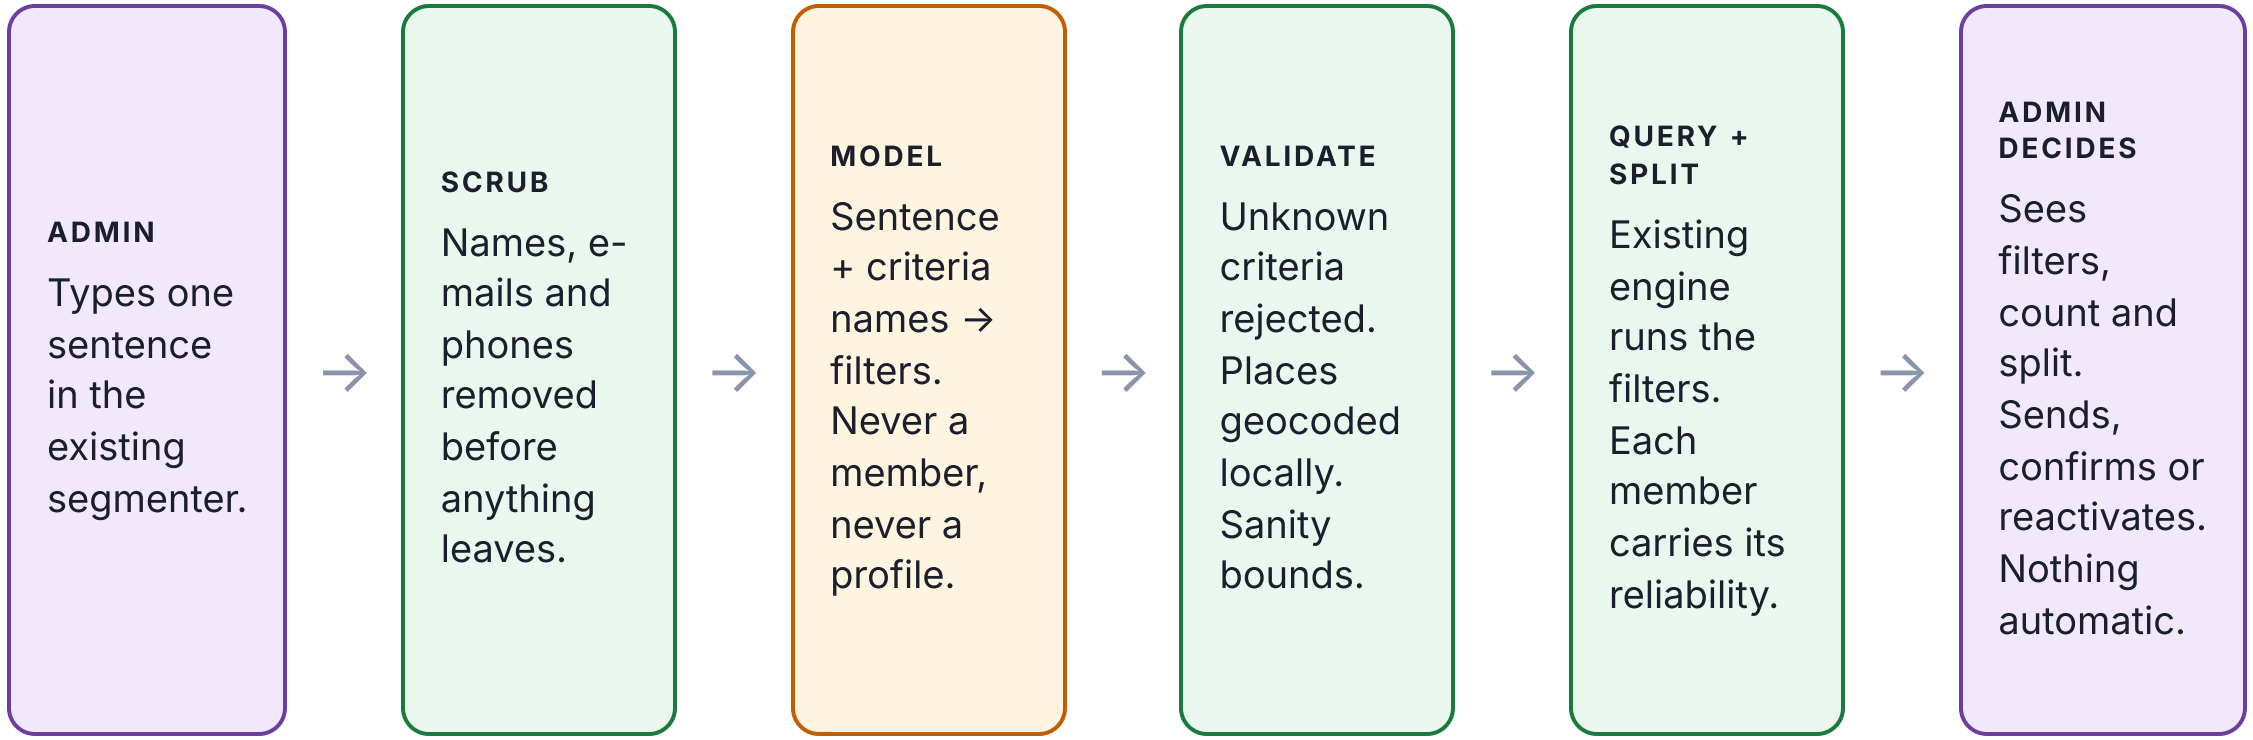
Current · To confirm · To reactivate

Three words, three actions: send, confirm, reactivate. Members never see a label, only "Profile last updated 4 years ago · Confirm / Correct".

Of the 49 matches: 1 current · 5 to confirm · 43 to reactivate

HOW IT WORKS

One model call, everything else deterministic



For the DPO, one sentence: only the administrator's scrubbed question and the names of the filters leave the platform. Never a member's profile.

Flag off = segmenter identical to today. Tenants who refuse any transfer keep everything except the sentence box.

THE FIRST INCREMENT · ONE QUARTER

Reliability layer first, sentence box second

WEEKS 1–2

Baselines

Instrument today's segmenter. Golden set of 100 sentences. Thresholds hand-checked by CS. Legal file opened.

WEEKS 3–7

Reliability layer

Nightly reliability tag, two new criteria on native date filters, reference tables, flags and metering. No engine refactor.

WEEKS 6–9

Sentence box

Gateway, filter contract, validator, pre-fill. Eval gate at 85 % before any tenant sees it.

WEEKS 8–12

Pilot, 10 tenants

Reliability from week 8, sentence box from week 10. Expert, mid and small tenants, one refusing transfer.

WEEKS 12 · 17

Two decisions

Week 12: keep or stop the sentence box. Week 17: keep or stop the reactivation loop.

26 developer-weeks of core, planned on roughly half of 3–4 shared developers (24–28 effective). Duplicate report and relative-date operator are already Q2. The reliability tag never moves.

Not in Q1: free-text chat, RAG over profiles, auto-correction, external enrichment, merge UI. Each with a reason in the full deck.

MEASURE, MONEY, RISK

How we know it works, and what pays for it

Stop criteria

Sentence box stops at week 12 if under 30 % of new segments use it or admins rewrite more than half of what it produces.

Reactivation stops at week 17 if fewer than 10 % of "to reactivate" profiles come back.

Money

Reliability included for everyone: the renewal argument for 350 tenants.

Sentence box included with fair use. Only the reactivation loop is priced, because only it is tied to dues collected.

Three risks, three positions

Dirty data: the reliability layer *is* the product.

DPO veto: no member data in any prompt, ever.

"We wanted ChatGPT": show the filters, end every demo on "6 you can e-mail".

Every threshold is an initial guess, reset on the week-4 baseline. Payback is formulas on measured baselines, not forecasts.

THE DECISION

Fund one quarter plus four weeks of measurement for Trusted Segments

"The administrator's problem is not building the list. It is believing it. So the first thing the platform returns is how far to trust it."

Full reasoning, exclusions, orchestration, legal file, SQL and operations: deck.pdf, 15 slides and four annexes.

Debanjan Mazumdar · Senior Product Manager Data & AI