ALUMNFORCE · SENIOR PM DATA & AI · DEBANJAN MAZUMDAR · OCTOBER 2026

# Trusted Segments

Describe your audience in one sentence. Get a segment you can trust. No member record ever leaves the platform.

**You asked for speed. Your users asked for trust.**
**So the platform returns not just a list, but how far to trust it.**



Full case study, annexes and SQL: deck.pdf

THE PROBLEM, REFORMULATED

# Eight of twelve verbatims are not about speed

**3**

**want speed**
"Saying it takes five seconds, doing it takes an hour."

**3**

**want a true list**
"340 London expats. Twenty had moved back years ago. Two unsubscribed."

**3**

**want proof**
"'Roughly' does not pass a board."

**2**

**want a data answer**
"Tell me exactly what leaves the platform."

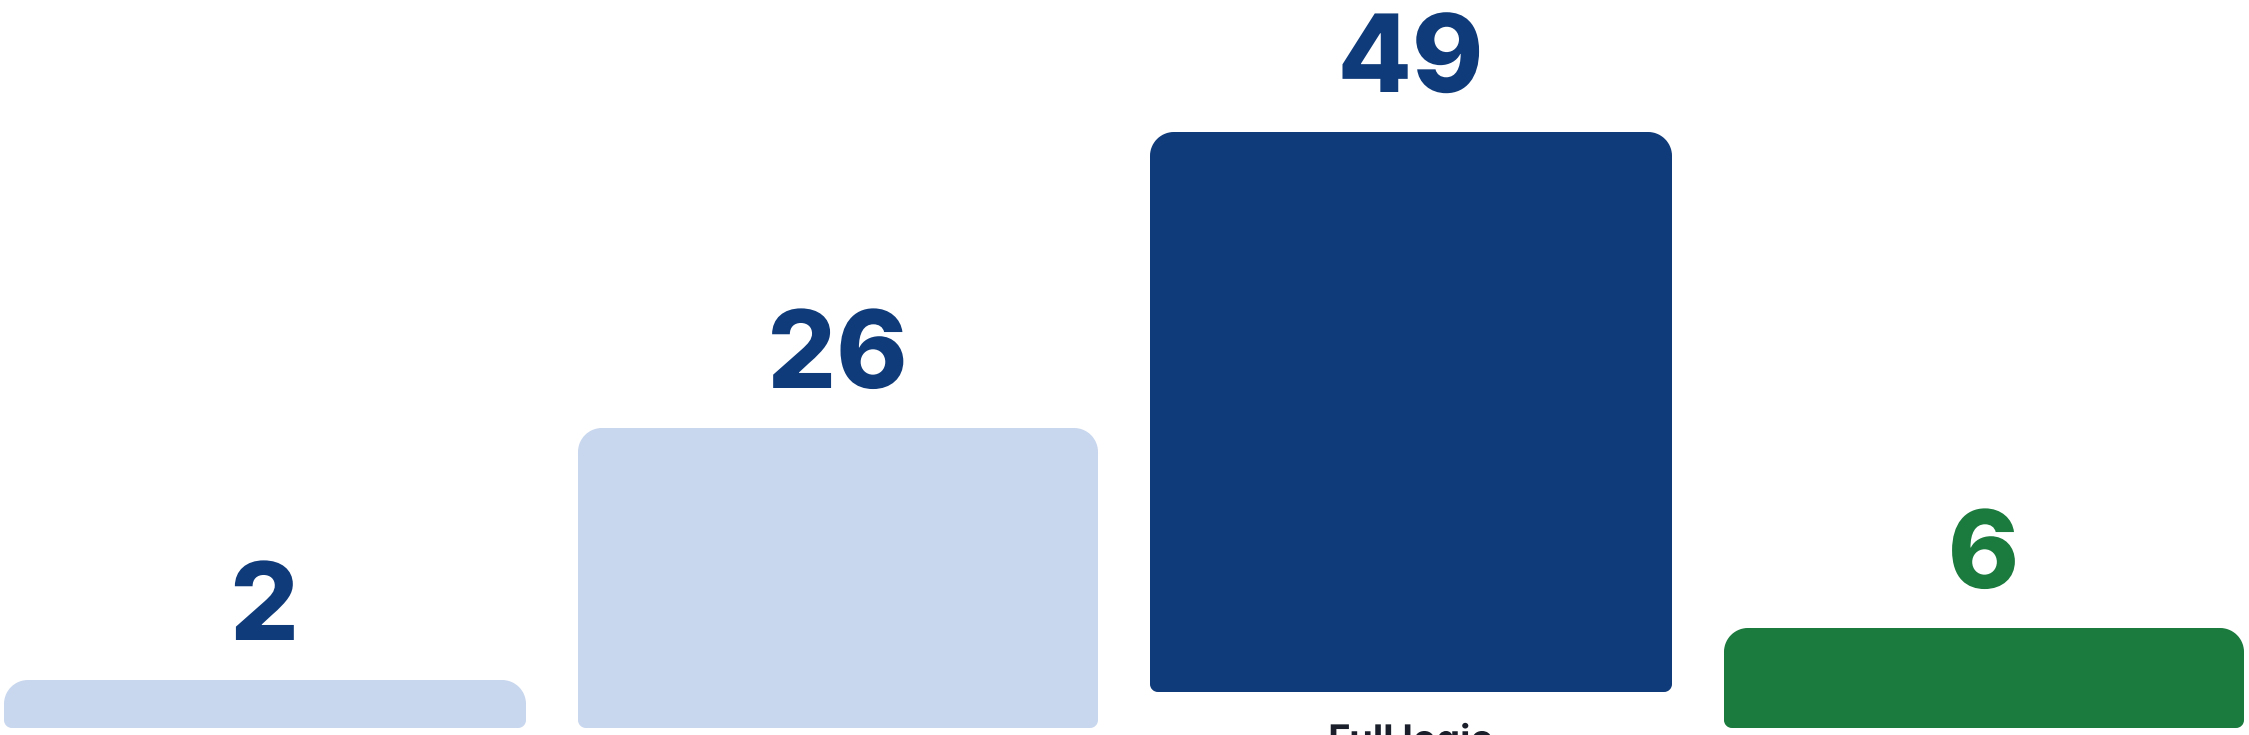
A chatbot that turns a sentence into filters serves the first three people. It makes the **wrong list faster** for everyone else.



One member says "not for me": 400 people, knows everyone, wants the Excel columns in order.

# Same sentence, four answers

"Graduates working in finance in London for more than three years"




**Naive filter**

city = "Londres", sector = "Finance", exact text

**Normalised**

8 spellings of London and 3 sector labels merged

**Full logic**

duplicates removed, employer and job title used as finance signals

**Worth e-mailing today**

profile or login recent enough to trust



68 % of profiles untouched for 2+ years · 525 rows for 500 announced members · 8 spellings of London

THE INSIGHT

# In this data, "three years in post" and "three years without an update" are the same signal

People update their profile when they start a job (median 112 days later) and then go silent. A naive assistant answering the brief's own example returns exactly the people **most likely to have moved**.

That is the 340-person London e-mail disaster, reproduced on their own file.

**So every segment gets a reliability split**

Current

To confirm

To reactivate

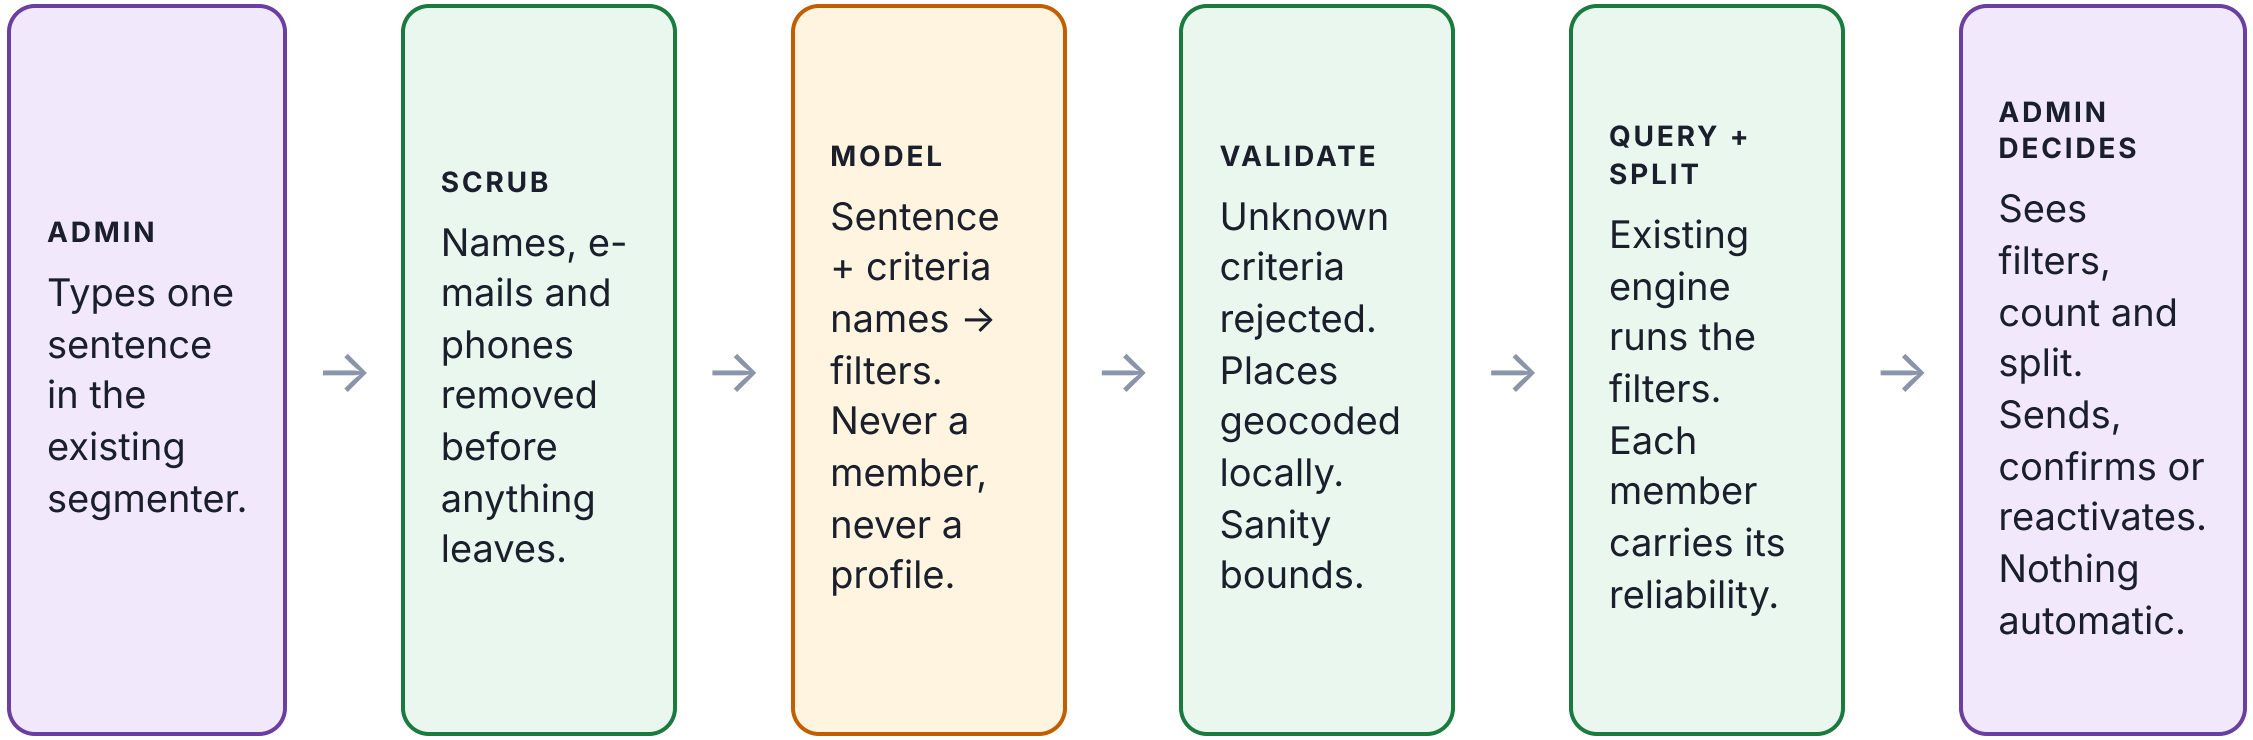
Three words, three actions: send, confirm, reactivate. Members never see a label, only "Profile last updated 4 years ago · Confirm / Correct".



Of the 49 matches: 1 current · 5 to confirm · 43 to reactivate

HOW IT WORKS

# One model call, everything else deterministic




**For the DPO, one sentence: only the administrator's scrubbed question and the names of the filters leave the platform. Never a member's profile.**



Flag off = segmenter identical to today. Tenants who refuse any transfer keep everything except the sentence box.

THE FIRST INCREMENT · ONE QUARTER

# Reliability layer first, sentence box second

WEEKS 1–2

## Baselines

Instrument today's segmenter. Golden set of 100 sentences. Thresholds hand-checked by CS. Legal file opened.

WEEKS 3–7

## Reliability layer

Status per member, two new criteria, normalised reference tables, flags and metering. The heaviest and most valuable part.

WEEKS 6–9

## Sentence box

Gateway, filter contract, validator, pre-fill. Eval gate at 85 % before any tenant sees it.

WEEKS 8–12

## Pilot, 10 tenants

Reliability from week 8, sentence box from week 10. Expert, mid and small tenants, one refusing transfer.

WEEKS 12 · 17

## Two decisions

Week 12: keep or stop the sentence box. Week 17: keep or stop the reactivation loop.

36 developer-weeks of core on 3 shared developers. If capacity drops, the duplicate report goes first. The reliability criteria never move.



Not in Q1: free-text chat, RAG over profiles, auto-correction, external enrichment, merge UI. Each with a reason in the full deck.

MEASURE, MONEY, RISK

# How we know it works, and what pays for it

### Stop criteria

Sentence box stops at week 12 if under 30 % of new segments use it or admins rewrite more than half of what it produces.

Reactivation stops at week 17 if fewer than 10 % of "to reactivate" profiles come back.

### Money

Reliability included for everyone: the renewal argument for 350 tenants.

Sentence box included with fair use. Only the reactivation loop is priced, because only it is tied to dues collected.

### Three risks, three positions

Dirty data: the reliability layer *is* the product.

DPO veto: no member data in any prompt, ever.

"We wanted ChatGPT": show the filters, end every demo on "6 you can e-mail".



Every threshold is an initial guess, reset on the week-4 baseline. Payback is formulas on measured baselines, not forecasts.

THE DECISION

# Fund one quarter plus four weeks of measurement for Trusted Segments

**"The administrator's problem is not building the list. It is believing it. So the first thing the platform returns is how far to trust it."**

Full reasoning, exclusions, orchestration, legal file, SQL and operations: deck.pdf, 15 slides and four annexes.



Debanjan Mazumdar · Senior Product Manager Data & AI